

Homework 2

Network Architecture I, Spring 2010

Due: paper submission, February 25th 9:30 am

(electronic submission on Bb and paper submission right before class)

1. Suppose within your Web browser you click on a link to obtain a web page. The IP address for the associated URL is cached in your local host, so a DNS look-up is not necessary to obtain the IP address. Further suppose that the Web page associated with the link references five very small objects on the same server. Let RTT_0 denote the RTTs between the local host and one of the object. Assuming zero transmission time of the object, how much time elapses from when the client clicks on the link until the client receives the full web page with
 - a. Nonpersistent HTTP?
 - b. Persistent HTTP?
2. Describe in detail how to register domain names and ip address for your start-up company (what companies to contact, how much is the fee, etc.).
3. What is an overlay network? Does it include routers? What are the edges in the overlay network? How is the query-flooding overlay network created and maintained?
4. What is meant by a handshaking protocol? What is/are an example(s) of handshaking protocol?

Laboratory Homework

Part 1: Telnet experiments

1. Try HTTP request (GET, HEAD, or POST) without using a web-browser. You can do this on command line using '> telnet *webserver* 80'. (for example, www.umkc.edu) Record the HTTP responses from the server – retrieve at least two different response status from the server.
2. Try sending an email without email composer. You can do this on command line using '> telnet *email_server* 25'. Send an email to yourself with a subject (title) and a short content. Then check the email with your regular email reader (Outlook, Eudora, etc.). Record your SMTP message as well as your received email.

Part 2: Wireshark experiments

In this part of the homework, you will use Wireshark to investigate HTTP protocol in operation. In this lab, you'll explore a couple of aspects of the HTTP protocol: the basic GET/response interaction, and retrieving large HTML files.

Download a packet trace file ([http-ethereal-trace-1](#)) of this homework from the Blackboard, Assignment Section. Or create a trace file on your own following the instruction at the footnote¹ from you home (not at University!).

Then, open the trace file on the Wireshark, and answer to the questions below.

Part2-1: The Basic HTTP GET/response interaction

Once you open the trace file, it will show in the packet-listing window that two HTTP messages were captured: the GET message (from your browser to the [gaia.cs.umass.edu](#) web server) and the response message from the server to your browser. The packet-contents window shows details of the selected message (in this case the HTTP GET message, which is highlighted in the packet-listing window). Recall that since the HTTP message was carried inside a TCP segment, which was carried inside an IP datagram, which was carried within an Ethernet frame, Wireshark displays the Frame, Ethernet, IP, and TCP packet information as well. We want to minimize the amount of non-HTTP data displayed (we're interested in HTTP here, and will be investigating these other protocols is later labs), so make sure the boxes at the far left of the Frame, Ethernet, IP and TCP information have a plus sign (which means there is hidden, undisplayed information), and the HTTP line has a minus sign (which means that all information about the HTTPmessage is displayed).

¹ FYI, below are the steps taken for the packet capture in the given trace. You don't collect packets yourself at school due to possible University policy issues.

1. Start up your web browser.
2. Start up the Wireshark packet sniffer. Enter "http" (just the letters, not the quotation marks) in the display-filter-specification window, so that only captured HTTP messages will be displayed later in the packet-listing window. (We're only interested in the HTTP protocol here, and don't want to see the clutter of all captured packets).
3. Wait a bit more than one minute (we'll see why shortly), and then begin Wireshark packet capture.
4. Enter the following to your browser
<http://gaia.cs.umass.edu/wireshark-labs/HTTP-wireshark-file1.html>
 Your browser should display the very simple, one-line HTML file.
5. Stop Wireshark packet capture.

(*Note:* You should ignore any HTTP GET and response for favicon.ico. If you see a reference to this file, it is your browser automatically asking the server if it (the server) has a small icon file that should be displayed next to the displayed URL in your browser. We'll ignore references to this pesky file in this lab.).

By looking at the information in the HTTP GET and response messages, answer the following questions. When answering the following questions, you should print out the GET and response messages (see the introductory Wireshark lab for an explanation of how to do this) and indicate where in the message you've found the information that answers the following questions.

1. Is your browser running HTTP version 1.0 or 1.1? What version of HTTP is the server running?
2. What languages (if any) does your browser indicate that it can accept to the server?
3. What is the IP address of your computer? Of the gaia.cs.umass.edu server?
4. What is the status code returned from the server to your browser?
5. When was the HTML file that you are retrieving last modified at the server?
6. How many bytes of content are being returned to your browser?
7. By inspecting the raw data in the packet content window, do you see any headers within the data that are not displayed in the packet-listing window? If so, name one.

(In your answer to question 5 above, if you have created your own trace file, you might have been surprised to find that the document you just retrieved was last modified within a minute before you downloaded the document. That's because (for this particular file), the gaia.cs.umass.edu server is setting the file's last-modified time to be the current time, and is doing so once per minute. Thus, if you wait a minute between accesses, the file will appear to have been recently modified, and hence your browser will download a "new" copy of the document.)

Part 2-2: Retrieving Long Documents

In the previous example, the document retrieved has been simple and short HTML files. Let's next see what happens when we download a long HTML file.

Download a packet trace file ([http-ethereal-trace-3](#)) of this homework from the Blackboard, Assignment Section. Or create a tracefile following the instruction at the footnote² from you home (not at University!)

Then, open the trace file on the Wireshark and answer to the questions below.

In the packet-listing window, you should see your HTTP GET message, followed by a multiple-packet response to your HTTP GET request. This multiple-packet response deserves a bit of explanation. Recall that the HTTP response message consists of a status line, followed by header lines, followed by a blank line, followed by the entity body. In the case of our HTTP GET, the entity body in the response is the *entire* requested HTML file. In our case here, the HTML file is rather long, and at 4500 bytes is too large to fit in one TCP packet. The single HTTP response message is thus broken into several pieces by TCP, with each piece being contained within a separate TCP segment (see Figure 1.24 in the text). Each TCP segment is recorded as a separate packet by Wireshark, and the fact that the single HTTP response was fragmented across multiple TCP packets is indicated by the "Continuation" phrase displayed by Wireshark. We stress here that there is no "Continuation" message in HTTP!

Answer the following questions:

1. How many HTTP GET request messages were sent by your browser?
2. How many data-containing TCP segments were needed to carry the single HTTP response?
3. What is the status code and phrase associated with the response to the HTTP GET request?
4. Are there any HTTP status lines in the transmitted data associated with a TCP-induced "Continuation"?

² Start up your web browser, and make sure your browser's cache is cleared (To do this under Firefox, select *Tools->Clear Private Data*, or for Internet Explorer, select *Tools->Internet Options->Delete File*; these actions will remove cached files from your browser's cache.)

1. Start up the Wireshark packet sniffer
2. Enter the following URL into your browser
<http://gaia.cs.umass.edu/wireshark-labs/HTTP-wireshark-file3.html>
 Your browser should display the rather lengthy US Bill of Rights.
3. Stop Wireshark packet capture, and enter "http" in the display-filter-specification window, so that only captured HTTP messages will be displayed.